

The road to reproducibility in Python

An introduction to the uv tool



Gaétan Lepage

November 15, 2025

GRICAD, Grenoble

Who am I?

Gaétan Lepage

- Ensimag 2020
- PhD @Inria Grenoble (RobotLearn team)
 - Deep Learning for robotic acoustics
- Nix(OS) contributor since 2021
 - Python ecosystem maintenance
 - Member of the CUDA team



Objectives

Objectives:

- Why “packaging” is important, especially in research
- Overview of how things work in Python
- A presentation of **uv**
- Hands-on! Try **uv on your own project**

Contents

| | |
|---|----|
| Reproducibility | 4 |
| Python packaging | 9 |
| uv: A modern approach on python tooling | 16 |
| Hands-on! | 33 |
| Conclusion | 35 |

Contents

| | |
|---|----|
| Reproducibility | 4 |
| Python packaging | 9 |
| uv: A modern approach on python tooling | 16 |
| Hands-on! | 33 |
| Conclusion | 35 |

Contents



Definitions: [1]

- **Repeatability:** Same team, same experimental setup
- **Reproducibility:** Different teams, same experimental setup
- **Replicability:** Different teams, different experimental setups

Contents



Definitions: [1]

- **Repeatability:** Same team, same experimental setup
- **Reproducibility:** Different teams, same experimental setup
- **Replicability:** Different teams, different experimental setups

Reproducibility crisis:

Difficulty to reproduce scientific studies from other groups

Contents



Definitions: [1]

- **Repeatability:** Same team, same experimental setup
- **Reproducibility:** Different teams, same experimental setup
- **Replicability:** Different teams, different experimental setups

Reproducibility crisis:

Difficulty to reproduce scientific studies from other groups

Only in experimental sciences... right?

Contents



Definitions: [1]

- **Repeatability:** Same team, same experimental setup
- **Reproducibility:** Different teams, same experimental setup
- **Replicability:** Different teams, different experimental setups

Reproducibility crisis:

Difficulty to reproduce scientific studies from other groups

Only in experimental sciences... right?

-> **Major issue in modern computer science research**

Reasons for the lack of reproducibility:

Reasons for the lack of reproducibility:

- Code is not available

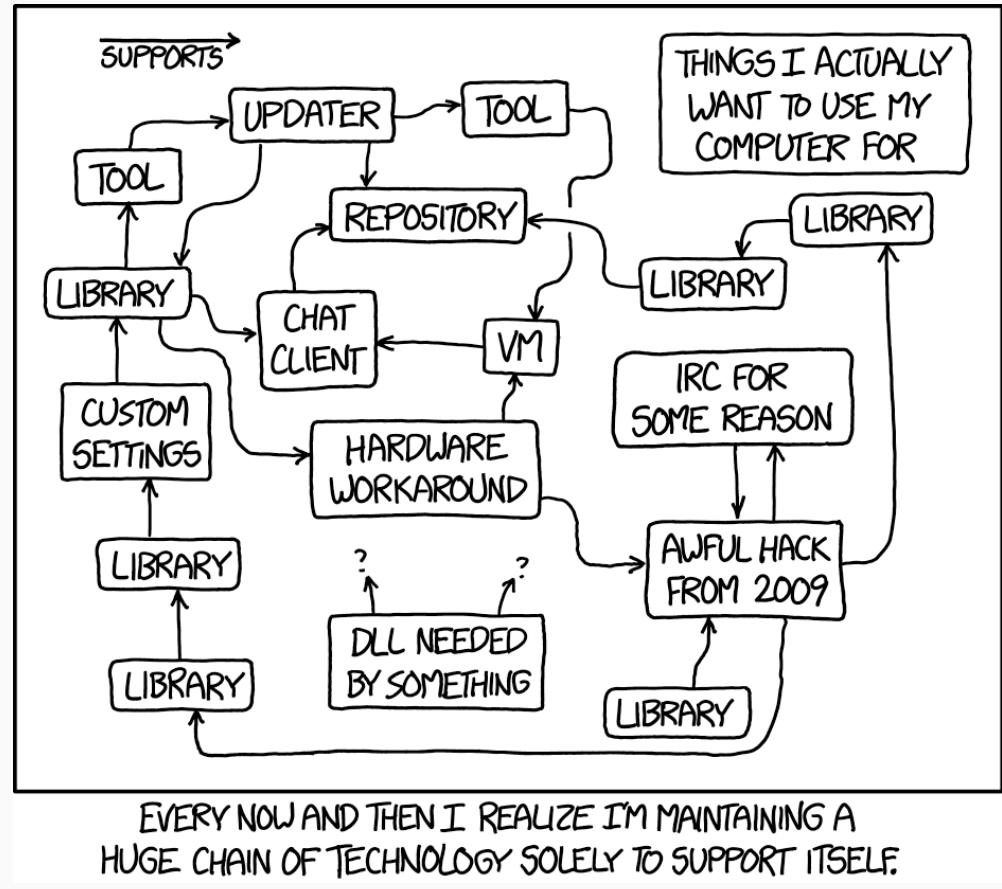
Reasons for the lack of reproducibility:

- Code is not available
- Data is not available

Software

Reasons for the lack of reproducibility:

- Code is not available
 - Data is not available
 - **Code is hard to run**



General recommendations

- Most importantly: keep things simple!
 - > Fewer dependencies/languages/constraints = fewer problems

General recommendations

- Most importantly: keep things simple!
 - > Fewer dependencies/languages/constraints = fewer problems
- Use version control (**git**)

General recommendations

- Most importantly: keep things simple!
 - > Fewer dependencies/languages/constraints = fewer problems
- Use version control ([git](#))
- Add a license ([LICENSE](#)) (MIT, GPL, Apache...)

General recommendations

- Most importantly: keep things simple!
 - > Fewer dependencies/languages/constraints = fewer problems
- Use version control ([git](#))
- Add a license ([LICENSE](#)) (MIT, GPL, Apache...)
- Write a [README.md](#) file:
 - Context about the project
 - Installation instructions
 - How to run the code?
 - How to download data?
 - How to replicate the results?

General recommendations

- Most importantly: keep things simple!
 - > Fewer dependencies/languages/constraints = fewer problems
- Use version control ([git](#))
- Add a license ([LICENSE](#)) (MIT, GPL, Apache...)
- Write a [README.md](#) file:
 - Context about the project
 - Installation instructions
 - How to run the code?
 - How to download data?
 - How to replicate the results?
- Ensure your code and experiments can be easily run and are reproducible

Existing approaches

- **Nothing**
 - > Fine for very simple software stacks. Doesn't scale

Existing approaches

- **Nothing**
 - > Fine for very simple software stacks. Doesn't scale
- **Natural language** (i.e. instructions)
 - > Documentation is great, but does not scale with complexity

Existing approaches

- **Nothing**
 - > Fine for very simple software stacks. Doesn't scale
- **Natural language** (i.e. instructions)
 - > Documentation is great, but does not scale with complexity
- **Virtual environments** (for Python)
 - > How to fill it? (not often reproducible)
 - > no control on the system's environment

Existing approaches

- **Nothing**
 - > Fine for very simple software stacks. Doesn't scale
- **Natural language** (i.e. instructions)
 - > Documentation is great, but does not scale with complexity
- **Virtual environments** (for Python)
 - > How to fill it? (not often reproducible)
 - > no control on the system's environment
- **What about containers?**
 - > How to fill it? (not often reproducible)
 - > Not resource-efficient

Existing approaches

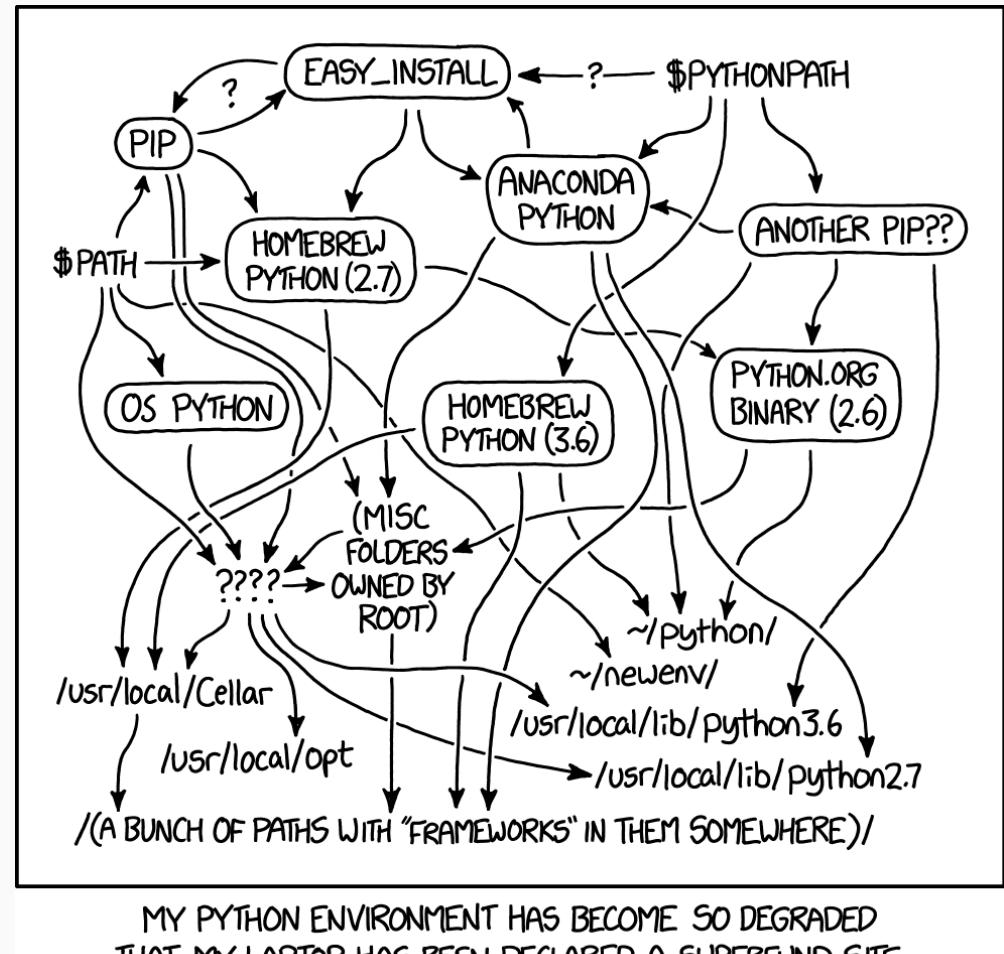
- **Nothing**
 - > Fine for very simple software stacks. Doesn't scale
- **Natural language** (i.e. instructions)
 - > Documentation is great, but does not scale with complexity
- **Virtual environments** (for Python)
 - > How to fill it? (not often reproducible)
 - > no control on the system's environment
- **What about containers?**
 - > How to fill it? (not often reproducible)
 - > Not resource-efficient
- **Functional package managers (Nix, Guix)**
 - > Elegant, powerful, but very hard to use

Contents

| | |
|---|----|
| Reproducibility | 4 |
| Python packaging | 9 |
| uv: A modern approach on python tooling | 16 |
| Hands-on! | 33 |
| Conclusion | 35 |

Contents

- Python projects often depend on many libraries
 - Many tools (*setuptools*, *pip*, *conda*, *poetry*)
 - Many standards (*setup.py*, *pyproject.toml*, *requirements.txt*, *conda-env.yml*)
- > Often laborious to “deploy” a project



Main components

Main components:

- The **build-system**: Runs when you invoke `pip install`
Ex: `setuptools`, `poetry`, `hatchling`, `uv-build`, ...

Main components

Main components:

- The **build-system**: Runs when you invoke `pip install`
Ex: `setuptools`, `poetry`, `hatchling`, `uv-build`, ...
- The dependencies specification Ex: `pyproject.toml`'s **dependencies** list (recommended),
`requirements.txt`, `setup.py`'s **install_requires** (`setuptools` only)

Main components

Main components:

- The **build-system**: Runs when you invoke `pip install`
Ex: `setuptools`, `poetry`, `hatchling`, `uv-build`, ...
- The dependencies specification Ex: `pyproject.toml`'s **dependencies** list (recommended),
`requirements.txt`, `setup.py`'s **install_requires** (`setuptools` only)
- Optionally, a project/environment management tool:
Ex: `uv`, `poetry`, `conda`, ...

Main components

Main components:

- The **build-system**: Runs when you invoke `pip install`
Ex: `setuptools`, `poetry`, `hatchling`, `uv-build`, ...
- The dependencies specification Ex: `pyproject.toml`'s `dependencies` list (recommended),
`requirements.txt`, `setup.py`'s `install_requires` (`setuptools` only)
- Optionally, a project/environment management tool:
Ex: `uv`, `poetry`, `conda`, ...
- Packaging repository **Pypi**:
 - Source distributions (`sdist`)
 - Binary builds (`wheels`)

Dependency specification

Formalized in [PEP-508](#).

Examples:

- `numpy`
- `torch==2.9.1`
- `pandas>=2.0,<2.4`
- `ray[data]`
- `requests [security,tests] >2.8.1,=2.8.* ; python_version < "2.7"`

Explanation:

- Dependency name: `requests`
- Optional features: `[security,tests]`
- Version constraints: `>2.8.1,=2.8.*`
- Platform compatibility algebra: `python_version < "2.7"`

From uv2nix talk by @adisbladis [2]

requirements.txt

```
requests
colorama; platform_system == "Windows"
importlib; python_version
numpy
torch>=2.8.0
tqdm
git+ssh://git@github.com/echweb/echweb-utils.git
git+https://github.com/DavidDiazGuerra/gpuRIR
```

- Came from Pip
- List of PEP-508 strings
- Usually used with
`pip install -r requirements.txt`
- Often alongside `setup.py`

From uv2nix talk by @adisbladis [2]

setup.py

```
from distutils.core import setup

setup(
    name="my_project",
    version="0.1.0",
    description="My great project",
    long_description=open('README.md').read(),
    install_requires=[
        "numpy",
        "torch>=2.8.0"
    ],
    author="Gaétan Lepage",
    author_email="gaetan@glepage.com",
    url="https://my-project.sh",
    license="MIT",
)
```

- Originated with `distutils/`
`setuptools`
- Most popular `build-system`
- Project metadata as Python code
- Build with `python setup.py build`
- Develop with `pip install -e`
- Not a standard
- Can be used for complex building (e.g.
native code compilation)

From uv2nix talk by @adisbladis [2]

pyproject.toml

```
name = "my_project"
version = "0.1.0"
description = "My great project"
readme = "README.md"
license = "MIT"

requires-python = ">=3.9,<3.14"
dependencies = [
    "numpy",
    "torch>=2.8.0"
]

[build-system]
requires = ["setuptools", "setuptools-scm"]
build-backend =
"setuptools.build_meta"
```

- Standard way of specifying the project metadata
- Specification: [PEP-517](#) and [PEP-621](#)
- Contains the list of dependencies (no more [requirements.txt](#))

From uv2nix talk by @adisbladis [2]

Contents

| | |
|---|----|
| Reproducibility | 4 |
| Python packaging | 9 |
| uv: A modern approach on python tooling | 16 |
| Hands-on! | 33 |
| Conclusion | 35 |

What is it?



What is **uv**?

- A modern tool to manage a python development project:
 - Specify dependencies...
 - ...install them...
 - reproducibly!

Developed by [Astral](#), a company building open source python tooling.

What is it?



Installation

```
curl -LsSf https://astral.sh/uv/install.sh | sh
```

Resources

- *Documentation:* <https://docs.astral.sh/uv>
- *Github repo:* <https://github.com/astral-sh/uv>
- *Python packaging documentation:* <https://packaging.python.org>

What is it?

- **Pros:**
 - Modern (August 2024)
 - Good ergonomics, intuitive UI
 - Fast (written in Rust)
 - Respects Python standards (e.g. PEP 621)
 - User-local installation (no need for `sudo`)
- **Cons:**
 - Limited to Python (v.s. Pixi [3] or Nix/Guix)

How does it work?

- Creates and edits the `pyproject.toml` file
- Downloads and manages its own Python interpreters
- Creates and modifies the virtual environment
 - Computes the versions of all dependencies (SAT solver)
 - Transitive dependencies are pinned too!
 - Save the result in the `uv.lock` file
 - Installs all dependencies in the virtual environment

Dependency locking: the key to *true* reproducibility

- All dependency versions are saved to the `uv.lock` lockfile
 - Ensures the environment can be reproduced later (same versions)
 - Generate/update the lock file: `uv lock`
 - Update dependencies: `uv lock --upgrade`
- > `uv.lock` should be tracked by `git`

How does it compare to other tools

- **Barebone pip (+ venv)**
 - Python based
 - Manages environments
 - “Slow”
 - No locking
- **Poetry, pdm et al.**
 - Python based
 - “Slow”
 - Consistent locking!
- **Conda**
 - Slow
 - Installs (some) *system libraries*
 - No locking

From *uv2nix* talk by @adisbladis [2]

Installation

On Linux and MacOS

```
$ curl -LsSf https://astral.sh/uv/install.sh | sh
```

On Windows

```
powershell -ExecutionPolicy ByPass -c "irm https://astral.sh/uv/install.ps1 | iex"
```

No **sudo** privileges required

Test it!

```
$ uv --version
```

Project initialization (uv init)

-> Generates important files:

- `pyproject.toml` (if necessary)
- lockfile: `uv.lock`
- `.git` and `.gitignore`
- `main.py` (Hello-world example)

```
# Existing projects  
$ uv init .
```

```
# New project  
$ uv init new_project
```

```
$ tree -a -L1  
├── .git  
├── .gitignore  
├── main.py  
├── pyproject.toml  
├── .python-version  
└── README.md
```

```
# Create the venv and install python  
$ uv sync  
  
# Two new files  
$ tree -a -L1  
...  
├── uv.lock  
└── .venv
```

Note: To specify a python version:

```
$ uv init --python 3.12
```

Adding dependencies

Most important command: `uv add`. It is the best way to *add* a dependency to the project

What it does:

- Adds the dependency specification (`numpy>=2.0,<2.4.0`) to `dependencies` in `pyproject.toml`
- Creates the virtual environment (`.venv`) if necessary
- Solves the environment (choose the version to install for each dependency)
- Installs the required dependencies (including transitive ones) to the virtual environment
- Updates `uv.lock`

You can provide arbitrary dependency specifications:

```
# Add a single dependency
$ uv add "requests[security,tests]>2.8.1"
```

```
# Add all dependencies from `requirements.txt`
$ uv add -r requirements.txt
```

```
# Remove a dependency
$ uv remove requests
```

Running the code

Two solutions to run the code:

1. `uv run main.py`
 - Automatically and transparently activates the virtual environment on the fly
2. Manually activate the virtual environment

```
$ source .venv/bin/activate
```

```
$ python main.py
```

Useful to quickly run scripts or the REPL

Dependency groups and optional dependencies

- **optional-dependencies**: for extra features

- ▶ Adding an optional dependency

```
$ uv add matplotlib --optional plot
```

- ▶ In `pyproject.toml`

```
[project.optional-dependencies]
plot = [
    "matplotlib>=3.6.3"
]
excel = [
    "odfpy",
    "xlsxwriter>=3.0.5"
]
```

- ▶ Installing (in another project):

```
$ uv add my_project[plot]
$ pip install my_project[plot]
```

- **dependency-groups**: for development dependencies

- ▶ Adding a development dependency:

```
$ uv add --group test pytest
```

- ▶ In `pyproject.toml`

```
[dependency-groups]
dev = [
    "pytest"
]
lint = [
    "ruff"
]
```

<https://pydevtools.com/handbook/explanation/what-are-optional-dependencies-and-dependency-groups/>

The `uv pip` command

uv implements most (all?) **pip** commands:

Example:

```
$ uv pip install numpy
```

- Much faster than the original **pip**
- Can sometimes be useful, but should not be used to install dependencies
- Prefer **uv add**

Storage management

uv stores data in multiple places:

- **Cache:**
 - ▶ uv uses aggressive caching to avoid re-downloading (and re-building) dependencies that have already been accessed in prior runs.
 - ▶ Contains downloaded and built dependencies, then linked in the virtual environments.
 - ▶ Where ? `~/.cache/uv` (`--cache-dir`, `$UV_CACHE_DIR`, `$XDG_CACHE_HOME/uv`)
 - ▶ Can be purged: `uv cache clean`

Storage management

`uv` stores data in multiple places:

- **Cache:**
 - ▶ `uv` uses aggressive caching to avoid re-downloading (and re-building) dependencies that have already been accessed in prior runs.
 - ▶ Contains downloaded and built dependencies, then linked in the virtual environments.
 - ▶ Where ? `~/.cache/uv (--cache-dir, $UV_CACHE_DIR, $XDG_CACHE_HOME/uv)`
 - ▶ Can be purged: `uv cache clean`
- **Configuration:**
 - ▶ Where ? `~/.config/uv`

Storage management

`uv` stores data in multiple places:

- **Cache:**
 - ▶ `uv` uses aggressive caching to avoid re-downloading (and re-building) dependencies that have already been accessed in prior runs.
 - ▶ Contains downloaded and built dependencies, then linked in the virtual environments.
 - ▶ Where ? `~/.cache/uv` (`--cache-dir`, `$UV_CACHE_DIR`, `$XDG_CACHE_HOME/uv`)
 - ▶ Can be purged: `uv cache clean`
- **Configuration:**
 - ▶ Where ? `~/.config/uv`
- **Persistent data directory:**
 - ▶ Contains `python` interpreters and tools
 - ▶ `~/.local/share/uv` (`$XDG_DATA_HOME/uv`)

Storage management

`uv` stores data in multiple places:

- **Cache:**

- `uv` uses aggressive caching to avoid re-downloading (and re-building) dependencies that have already been accessed in prior runs.
- Contains downloaded and built dependencies, then linked in the virtual environments.
- Where ? `~/.cache/uv (--cache-dir, $UV_CACHE_DIR, $XDG_CACHE_HOME/uv)`
- Can be purged: `uv cache clean`

- **Configuration:**

- Where ? `~/.config/uv`

- **Persistent data directory:**

- Contains `python` interpreters and tools
- `~/.local/share/uv ($XDG_DATA_HOME/uv)`

- **Virtual environment:**

- By default, in `my_project/.venv/`
- Contains links to the cache directory (must be on the same FS)
- Contains all project dependencies

WARNING: Be careful when running on systems where `$HOME` storage is limited.

Use `uvx` to install/run an executable on the fly

```
# Run a tool, right here, right now  
$ uvx ruff  
$ uv tool run ruff # same, but more verbose  
  
# Add dependencies on the fly  
$ uvx --with pandas,pyarrow ipython  
  
$ uvx --from jupyter-core jupyter lab
```

To install a CLI tool with `uv`:

```
$ uv tool install ruff  
$ which ruff  
/home/gaetan/.local/bin/ruff
```

Deploying your projects to a remote system

Install `uv`, clone your project and run your code, that's it!

```
$(local) [~/work/project] git push
$(local) [~/work/project] ssh cluster

$(cluster) [~]           git clone <PROJECT_URL>
$(cluster) [~/project]   cd project
$(cluster) [~/project]   uv run main.py
```

uv in Docker containers

- `uv` can be used in Docker containers
- Both *distroless* and regular images are provided. `uv` is pre-installed
 - ▶ Distroless: ghcr.io/astral-sh/uv:latest
 - ▶ Alpine: ghcr.io/astral-sh/uv:alpine
 - ▶ Debian: ghcr.io/astral-sh/uv:debian-slim

uv in Docker containers

- uv can be used in Docker containers
 - Both *distroless* and regular images are provided. uv is pre-installed
 - Distroless: `ghcr.io/astral-sh/uv:latest`
 - Alpine: `ghcr.io/astral-sh/uv:alpine`
 - Debian: `ghcr.io/astral-sh/uv:debian-slim`
-

- Run your app!

```
FROM ghcr.io/astral-sh/uv:debian-slim
ENV UV_COMPILE_BYTECODE=1 UV_LINK_MODE=copy

# Copy the project into the image
ADD . /project

# Sync the project into a new environment, asserting the lockfile is up to date
WORKDIR /project

RUN uv sync --locked

# Presuming there is an `hello` command provided by the project
CMD ["uv", "run", "hello"]
```

<https://docs.astral.sh/uv/guides/integration/docker>

Contents

| | |
|---|----|
| Reproducibility | 4 |
| Python packaging | 9 |
| uv: A modern approach on python tooling | 16 |
| Hands-on! | 33 |
| Conclusion | 35 |

Contents

Installation

```
curl -LsSf https://astral.sh/uv/install.sh | sh
```

- *Documentation:* <https://docs.astral.sh/uv>
 - *GitHub repo:* <https://github.com/astral-sh/uv>
 - *Python packaging documentation:* <https://packaging.python.org>
-

Contents

Installation

```
curl -LsSf https://astral.sh/uv/install.sh | sh
```

- Documentation: <https://docs.astral.sh/uv>
 - GitHub repo: <https://github.com/astral-sh/uv>
 - Python packaging documentation: <https://packaging.python.org>
-

Your turn!

- Pick a project:
 - ▶ Your own Python project
 - ▶ One of your students/colleagues' project
 - ▶ Open source code from an article
- Bootstrap uv

Installation

```
curl -LsSf https://astral.sh/uv/install.sh | sh
```

- Documentation: <https://docs.astral.sh/uv>
 - GitHub repo: <https://github.com/astral-sh/uv>
 - Python packaging documentation: <https://packaging.python.org>
-

Your turn!

- Pick a project:
 - Your own Python project
 - One of your students/colleagues' project
 - Open source code from an article
- Bootstrap `uv`

Cheatsheet:

```
# Init a project  
$ uv init . # or uv init  
my_project
```

```
# Add a dependency  
$ uv add numpy
```

```
# Run the code  
$ uv run main.py
```

```
# Sync the virtual environment  
$ uv sync
```

Any questions?

Contact:

-  gaetan@glepage.com
-  <https://glepage.com>
-  [GaetanLepage](#)

Conclusion

- [1] B. Antunes and D. R. Hill, “Reproducibility, Replicability and Repeatability: A survey of reproducible research with a focus on high performance computing,” *Computer Science Review*, vol. 53, p. 100655, 2024, doi: <https://doi.org/10.1016/j.cosrev.2024.100655>.
- [2] adisbladis, “Python packaging with nixpkgs, pyproject.nix & uv2nix, NixCon 2025.” [Online]. Available: <https://talks.nixcon.org/nixcon-2025/talk/Y8TSAW/>
- [3] T. Fischer *et al.*, “Pixi: Unified Software Development and Distribution for Robotics and AI,” *arXiv preprint arXiv:2511.04827*, 2025.